

CS610 Assignment 2

1 Problem 1

Performance Bugs Identified:

False Sharing

The `struct tracker` has `word_count` array whose size is equal to the number of threads. In the execution of the statement `tracker.word_count[thread_id]++`; (Line 142 in `reference.cpp`), false sharing is present which is clear from the code as well as from the generated PERF report.

To remove this false sharing, padding was done in the `word_count` array and its size was made equal to $\text{NUM_THREADS} \times 8$ and instead of accessing the index `thread_id`, `thread_id*8` was accessed by thread having `tid` equal to `thread_id`. This multiplication by 8 ensured that the memory location accessed by each thread is on separate cache line (8 `uint_64` of 8 Bytes each take 64 Bytes, which is equal to the cache line size) and thus false sharing is minimised. The source file `padded.cpp` has these modifications.

True Sharing

Taking and releasing lock for every word and line of the input file caused true sharing and was the cause of large number of HITMs (Lines 144-146 and 133-135 in `reference.cpp`). This was reduced by keeping thread-private variables for the words and line counts of each thread, and then later reducing them with the count in `tracker` struct, thus taking lock only once. This reduced the HITMs to nearly zero. The source file `improved.cpp` has these modifications. Also we have removed the redundant mutex `line_count_mutex` since the mutex in `struct tracker` is sufficient for the synchronisation needed.

Shared Cache Line Distribution Pareto																
#	----- HITM -----		----- Store Refs -----		----- Data address -----		----- cycles -----		Total	cpu						
#	Num	RmtHitm	LclHitm	L1 Hit	L1 Miss	N/A	Offset	Node	PA cnt	Code address	rmt hitm	lcl hitm	load	records	cnt	Symbol
#	Object			Source:Line	Node{cpu list}											
#	0	0	3054	0	0	0	0x5d23dcb17380									
	0.00%	10.09%	0.00%	0.00%	0.00%	0{4,8-10,15}	0x0	0	1	0x5d23dcb134d1	0	190	169	473	5	[.] thread_runner(void*)
	reference.out	thread_runner(void*)+817	0{4,8-10,15}	0.00%	0.00%	0{4,8-10,15}	0x8	0	1	0x5d23dcb134d1	0	207	163	459	6	[.] thread_runner(void*)
	0.00%	9.43%	0.00%	0.00%	0.00%	0{0,4,6,8,10,12}	0x10	0	1	0x5d23dcb134d1	0	199	159	482	5	[.] thread_runner(void*)
	reference.out	thread_runner(void*)+817	0{0,4,6,8,10,12}	0.00%	0.00%	0{0,4,6,8,10,12}	0x18	0	1	0x5d23dcb134d1	0	214	178	468	6	[.] thread_runner(void*)
	0.00%	9.40%	0.00%	0.00%	0.00%	0{0,4,6,9-11}	0x20	0	1	0x5d23dcb134d1	0	190	161	434	4	[.] thread_runner(void*)
	reference.out	thread_runner(void*)+817	0{0,4,6,9-11}	0.00%	0.00%	0{0,4,6,9-11}	0x28	0	1	0x5d23dcb1332a	0	187	131	257	11	[.] thread_runner(void*)
	0.00%	8.35%	0.00%	0.00%	0.00%	0{0,3-4,6,8-12,14-15}	0x30	0	1	0x5d23dcb134de	0	225	154	1886	11	[.] thread_runner(void*)
	reference.out	thread_runner(void*)+830	0{0,3-4,6,8-12,14-15}	0.00%	0.00%	0{0,3-4,6,8-12,14-15}	0x38	0	1	0x75eaf4e99aa4	0	462	92	2905	11	[.] pthread_mutex_unlock@@GLIBC
	0.00%	16.21%	0.00%	0.00%	0.00%	0{0,3-4,6,8-12,14-15}	0x38	0	1	0x75eaf4e97f40	0	389	192	2746	11	[.] pthread_mutex_lock@@GLIBC
	libc.so.6	pthread_mutex_unlock.c:43	0{0,3-4,6,8-12,14-15}	0.00%	0.00%	0{0,3-4,6,8-12,14-15}	0x38	0	1	0x75eaf4e91294	0	263	182	331	11	[.] __GI___lll_lock_wait
	0.00%	13.52%	0.00%	0.00%	0.00%	0{0,3-4,6,8-12,14-15}	0x38	0	1	0x75eaf4e912a3	0	346	138	1716	11	[.] __GI___lll_lock_wait
	0.00%	5.27%	0.00%	0.00%	0.00%	0{0,3-4,6,8-12,14-15}	0x38	0	1							
	libc.so.6	lowlevellock.c:42	0{0,3-4,6,8-12,14-15}	0.00%	0.00%	0{0,3-4,6,8-12,14-15}	0x38	0	1							
	0.00%	3.18%	0.00%	0.00%	0.00%	0{0,3-4,6,8-12,14-15}	0x38	0	1							
	libc.so.6	lowlevellock.c:45	0{0,3-4,6,8-12,14-15}	0.00%	0.00%	0{0,3-4,6,8-12,14-15}	0x38	0	1							

Figure 1: PERF Report of Reference Version - Large number of HITMs in the `thread_runner` function

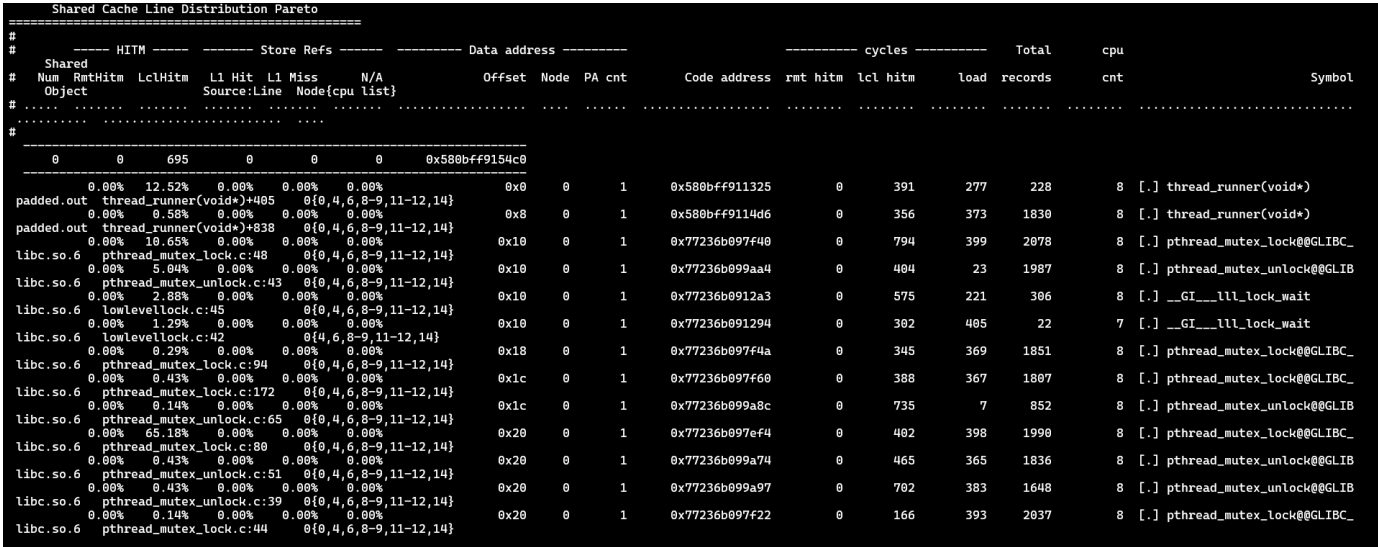


Figure 2: PERF Report of Padded Version - Small number of HITMs in the `thread_runner` function, majority of them are due to locking



Figure 3: PERF Report of Improved Version - No HITMs

Version	HITMs	Time (s)
Reference	5904	1.651
Padded	1291	1.009
Improved	0	0.070

Table 1: Performance Comparison of three versions (Data taken on `csews172` with 11 MB files)

Running the Code

There is a Makefile and a bash script in the code. `make` generates the executables for three versions. `run.sh` generates the PERF report for all three versions.

make

```
./run.sh <path to the file which has the names of input files>
```

The three PERF reports will be generated as `perf_report_reference.out`, `perf_report_padded.out` and `perf_report_improved.out`.

Also, there is a file `generate_input.cpp` which takes an integer `i` as CLI and generates a large random text file `inp<i>.txt` in the `test1` directory. `generate_input.out` is also generated by the `make` command. Ensure that `test1` directory is present before running `generate_input.out`.

2 Problem 2

The directory `problem2-dir` contains the source files. Below is the compilation and execution commands -

```
g++ problem2.cpp -lpthread
./a.out -inp=<input_path> -thr=<num_producers> -lns=<lines_per_thread> -buf=<shared_buffer_size>
-out=<output_path>
```

3 Problem 3

```
1  for i = 1 , N - 2
2      for j = i + 1 , N
3          A ( i , j - i ) = A ( i , j - i - 1 ) - A ( i + 1 , j - i ) + A ( i - 1 , i + j - 1 )
```

We consider the following pairs and type of dependencies:

1. $A(i, j - i)$ and $A(i, j - i - 1)$

$$i_0 = i_0 + \Delta i \implies \boxed{\Delta i = 0}$$

$$j_0 - i_0 = j_0 + \Delta j - i_0 - \Delta i - 1 \implies \Delta j = \Delta i + 1 \implies \boxed{\Delta j = 1}$$

Thus the flow dependency is $(0, 1)$ or $(0, +)$, which is a valid flow dependency. Since, we have a valid flow dependency, the corresponding anti dependency would be invalid.

2. $A(i, j - i)$ and $A(i + 1, j - i)$

$$i_0 = i_0 + \Delta i + 1 \implies \boxed{\Delta i = -1}$$

$$j_0 - i_0 = j_0 + \Delta j - i_0 - \Delta i \implies \Delta j = \Delta i \implies \boxed{\Delta j = -1}$$

The flow dependency is $(-1, -1)$, but it is an invalid dependency since the first non-zero component of the vector is negative. However, there will be an anti dependency with the vector equal to $(1, 1)$.

3. $A(i, j - i)$ and $A(i - 1, i + j - 1)$

$$i_0 = i_0 + \Delta i - 1 \implies \boxed{\Delta i = 1}$$

$$j_0 - i_0 = i_0 + \Delta i + j_0 + \Delta j - 1 \implies \boxed{\Delta j = -2i_0}$$

In this case, Δi is positive and $\Delta j = -2i_0$. Again since i_0 is always positive, Δj will be always negative. Thus the dependency vector will be $(+, -)$. It is a valid flow dependency and so there is no anti dependency.

There is no output dependency for any of the memory access pairs.